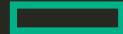


Container Orchestration: Which Conductor?

ContainerCon Europe, Berlin, Oct 2016



Hewlett Packard
Enterprise

Mike Bright, @mjbright



Haikel Guemar, @hguemar



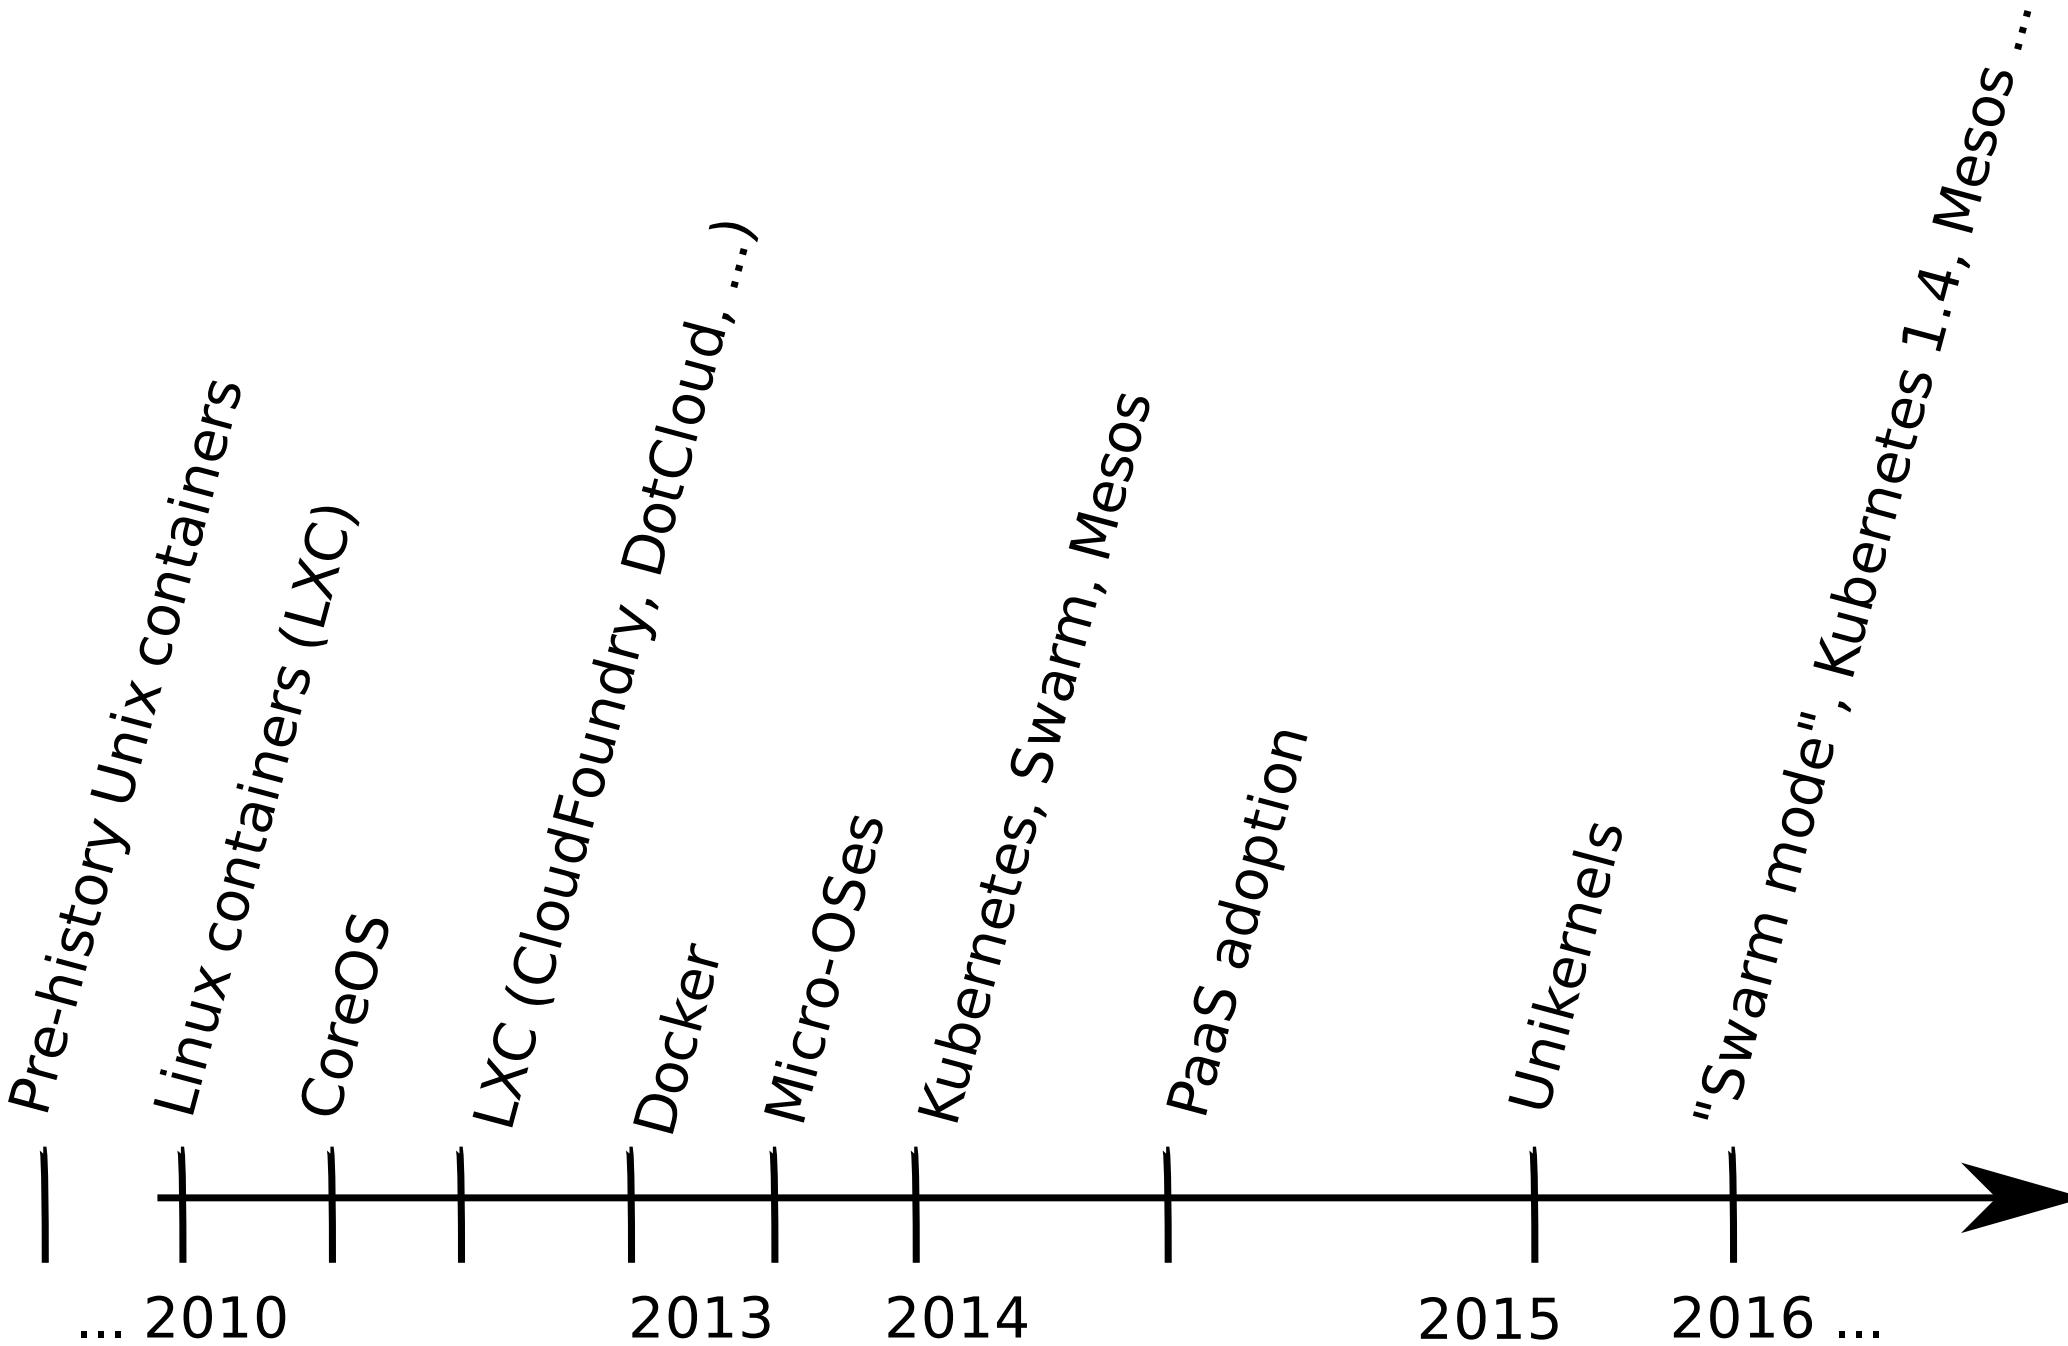
Mario Loriedo, @mariolet

First ...

First ...

A little bit of history

So let's first look at recent container history ...

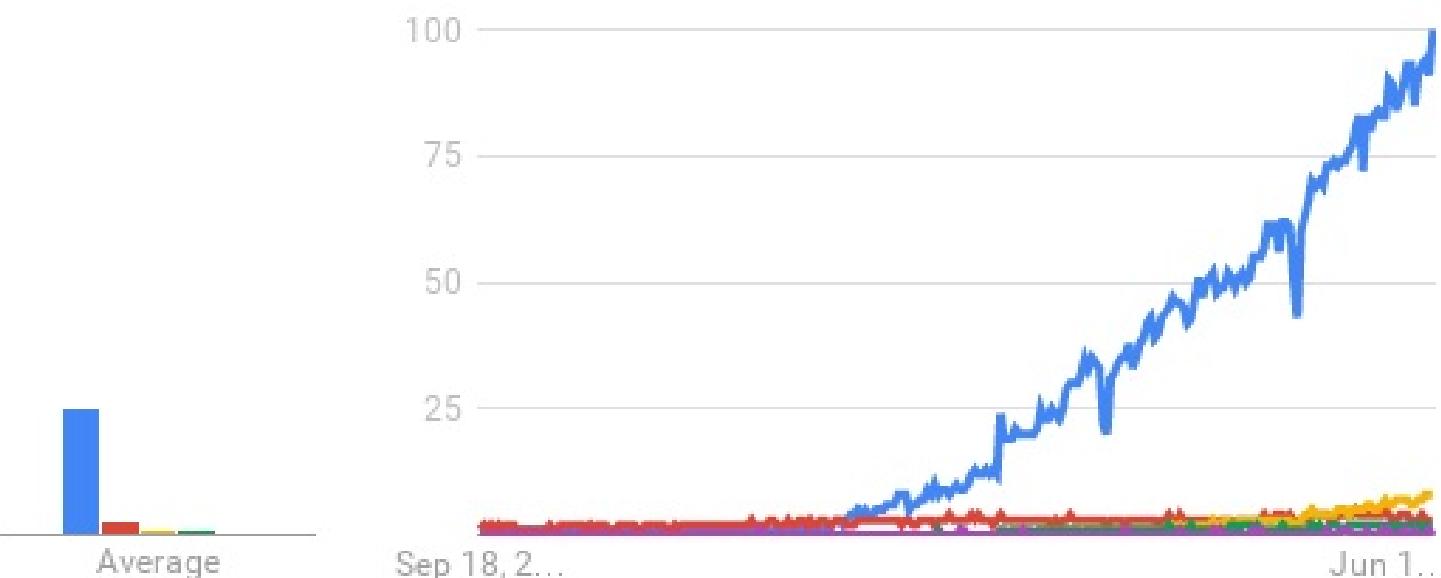


@hguemar @mjbright @mariolet

Interest over time

Google Trends

Docker lxc kubernetes Mesos Containerization



Worldwide. Past 5 years.

@hguemar @mjbright @mariolet

History μ-OSes

Many vendors are developing μ-OSes, small OS (mainly Linux-based) to be the basis for container engine hosts whether they be bare-metal or virtual **host machines**.

They're small, with fast startup, use few resources and have a small attack surface and often "*atomic*" software updates.

OS	Vendor
CoreOS	- (CoreOS)
Project Atomic	- (RedHat)
Photon	- (VMWare)
RancherOS	- (Rancher Labs)
Nano Server OS	- (Microsoft)
Ubuntu Snappy Core	- (Canonical)

History μ-OSes

Many vendors are developing μ-OSes, small OS (mainly Linux-based) to be the basis for container engine hosts whether they be bare-metal or virtual **host machines**.

They're small, with fast startup, use few resources and have a small attack surface and often "*atomic*" software updates.

OS	Vendor
CoreOS	- (CoreOS)
Project Atomic	- (RedHat)
Photon	- (VMWare)
RancherOS	- (Rancher Labs)
Nano Server OS	- (Microsoft)
Ubuntu Snappy Core	- (Canonical)

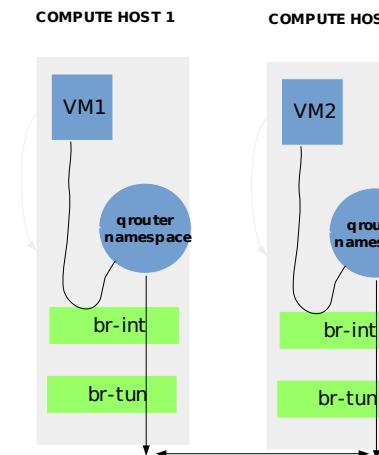
- ...Unikernels (...)

μ -Services

μ -services

From monoliths to μ -services

Remember when **high availability** meant this ...?



Servers running **monolithic applications** in **Active-Standby** modes, as 1+1, N+1, or N+M or split across 3 tiers.

Scaling meant to "**scale up**" by adding CPU, RAM, disk.
But there's a limit to this ... then you have to "**scale out**"

@bquemar @mibright @mariolnt

μ - services

From monoliths to μ - services

Then came μ -services ..

As the industry moved to virtualized micro-services this allowed to obtain greater efficiencies (higher utilisation of resources) and the redesign of applications allows to scale out and achieve high availability.

Containers facilitate this move, allowing faster scaling and even greater efficiencies with less redundancy (no OS to reproduce).

How containers help?

Container solutions such as Docker go beyond the isolation capabilities of LXC by providing simple to use tools to enable packaging of apps with their dependencies allowing portable applications between systems.

Containers are lightweight

Containers can be shared

Containers allow to use the same application binaries on development, test and production systems whether that be on a laptop, server or in the cloud.

μ -
services

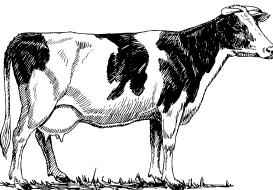
From monoliths to μ -services

But 1000's of nodes are unmanageable ... aren't they?

We can't take care of our



so we have to treat them like



that's cloud native !

@hguemar @mjbright @mariolet

So we need container orchestration



Orchestration

What was Container Orchestration again?

... and how does it differ from automation?

Orchestration is

- Architecture
 - Composition
 - Stitching
- Workflows
- Policies

Orchestration

We already have many choices for
Container Orchestration Engines

Docker Swarm : Docker Inc.

Kubernetes : Cloud Native Computing Foundation

Apache Mesos : Apache Software Foundation

Fleet : CoreOS

Rancher : Rancher Labs

Nomad : HashiCorp

Orchestration Imperative or Declarative

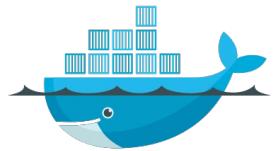
To manage 1000's of nodes it becomes necessary to be able to make declarative requests to the system.

	Imperative	Declarative
Tell the system	what to do <i>"start a new node"</i>	desired state <i>"3 mysql nodes"</i>
Intelligence	Operator	Orchestration Engine
Flexibility	Best	Least

Choice is great - when you know
what you want ...

Orchestration

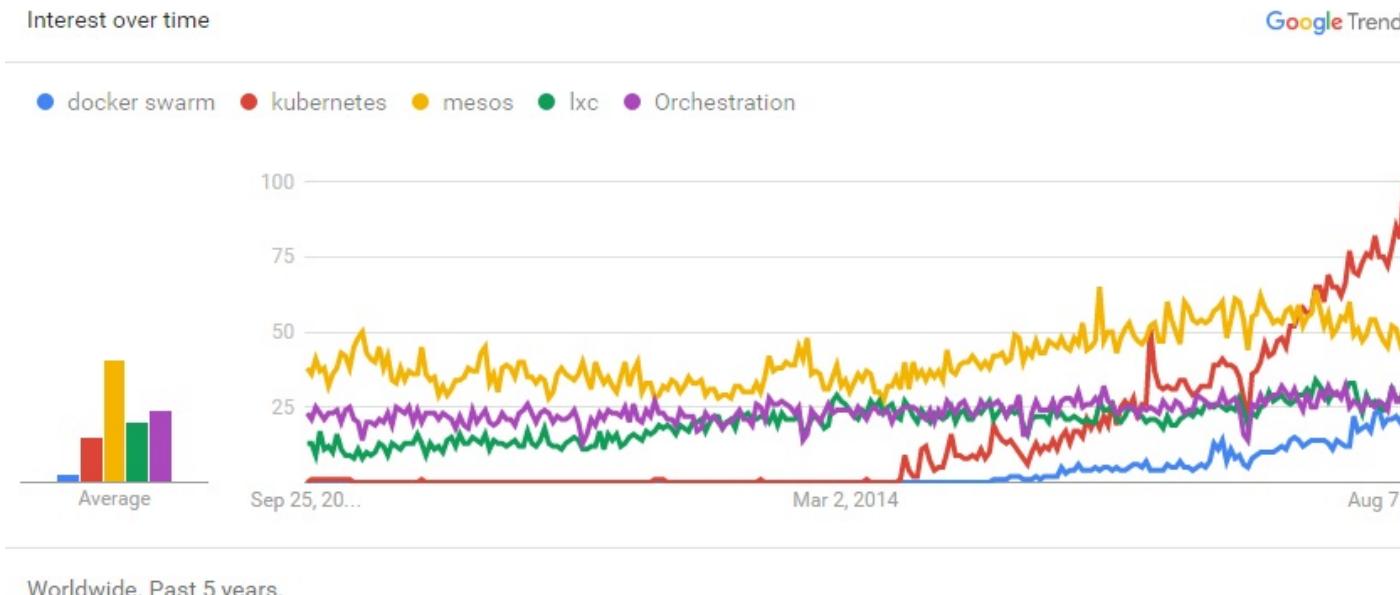
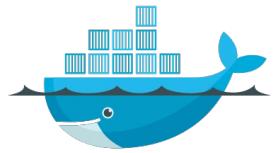
The Big 3 - Main Orchestration Choices



- Docker Swarm
 - Docker swarm
 - The swarm toolkit
 - Docker "**swarm mode**"
- Apache Mesos
 - Frameworks
 - Marathon (Mesosphere), Chronos, Aurora
 - Plugins
 - Jenkins
 - Mesosphere, DC/OS
- Kubernetes

@hguemar @mjbright @mariolet

Orchestration The Big 3 - What does Google Trends say?



Clearly Kubernetes has a lead, but we can expect *"Docker Swarm"* to quickly make progress thanks to the new *"swarm mode"*

But lets not forget the alternatives ...

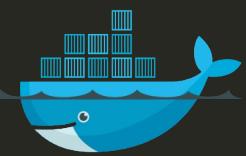
Orchestration More Choice ...



- Rancher (Rancher Labs)
- Fleet (CoreOS) A distributed init system (between systemd and etcd)
- Nomad (HashiCorp)

@hguemar @mjbright @mariolet

Docker Swarm





Docker Swarm

- 2014 Dec: Docker Swarm was announced
 - Capable of scaling up an architecture defined by Docker Compose
- 2015: Swarm Toolkit released
- 2015 June: Docker engine with Swarm Mode released with Docker 1.12

Docker "Swarm Mode" is quite a revolution in Docker Engine capabilities as it integrates

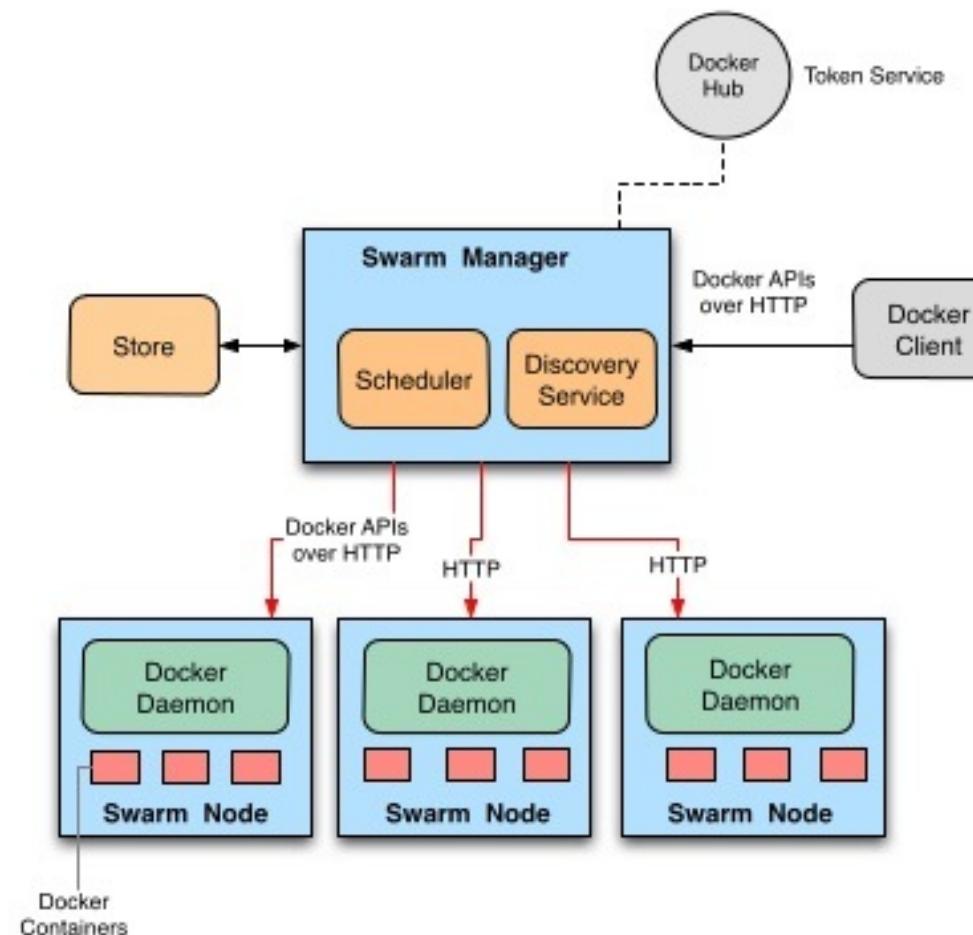
- Integration of Orchestration Capabilities directly within the Docker Engine
 - Load balancing across a mesh network
 - Service Discovery
- The same ease of use we're accustomed to from Docker !

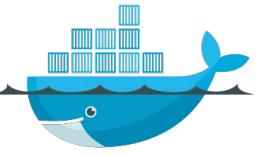


Architecture

Docker Swarm

Docker Swarm Architecture - Exploded





Docker Swarm

Using Docker "Swarm Mode"

Create a new swarm by creating the master node:

```
$ docker swarm init  
xxxx
```

Join a new node to the swarm swarm:

```
$ docker swarm join ....
```



Docker Swarm

Getting started

An excellent place to start is with Jerome Pettazoni's
"Orchestration Workshop"

- being run at this conference
- Available on github,
<https://github.com/jpetazzo/orchestration-workshop>

Apache Mesos





Apache Mesos

Arguably the most production ready orchestration today, exists since 2009.

Can scale to ~ 10,000 nodes.

Used in production by:

- Twitter
- Groupon
- Netflix

Mesos is used in conjunction with Frameworks such as

- **Marathon** or Aurora: manages long running tasks
- Chronos: designed for job orchestration
- Hadoop: for big data processing
- Kubernetes: allowing declarative use

Apache Mesos

Architecture

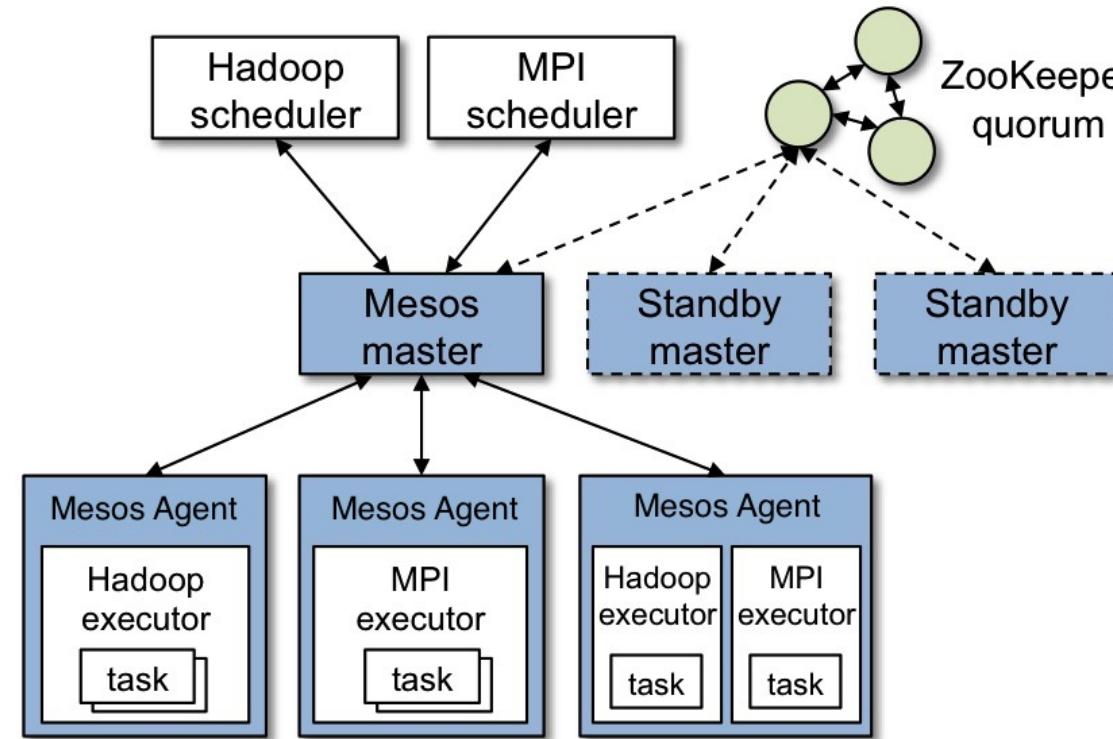


Image courtesy of
<http://mesos.apache.org/documentation/latest/architecture/>



Apache Mesos

Using Mesos

Step1

xxxx

Step2

Apache Mesos

Getting started

An excellent place to start is with the following tutorials

- Mesos
 - minimesos
- Mesosphere

Kubernetes



From the Greek: "Steersman, helmsman, sailing master"



Kubernetes is an open source project created by Google based on it's extensive experience running containers (millions of containers over a decade or so) from it's Borg and Omega projects.

Kubernetes

Started ~ Oct 2014, reach v1.0 in July 2015 and currently at v1.4 It is managed by the Cloud Native Computing Foundation <https://cncf.io/>

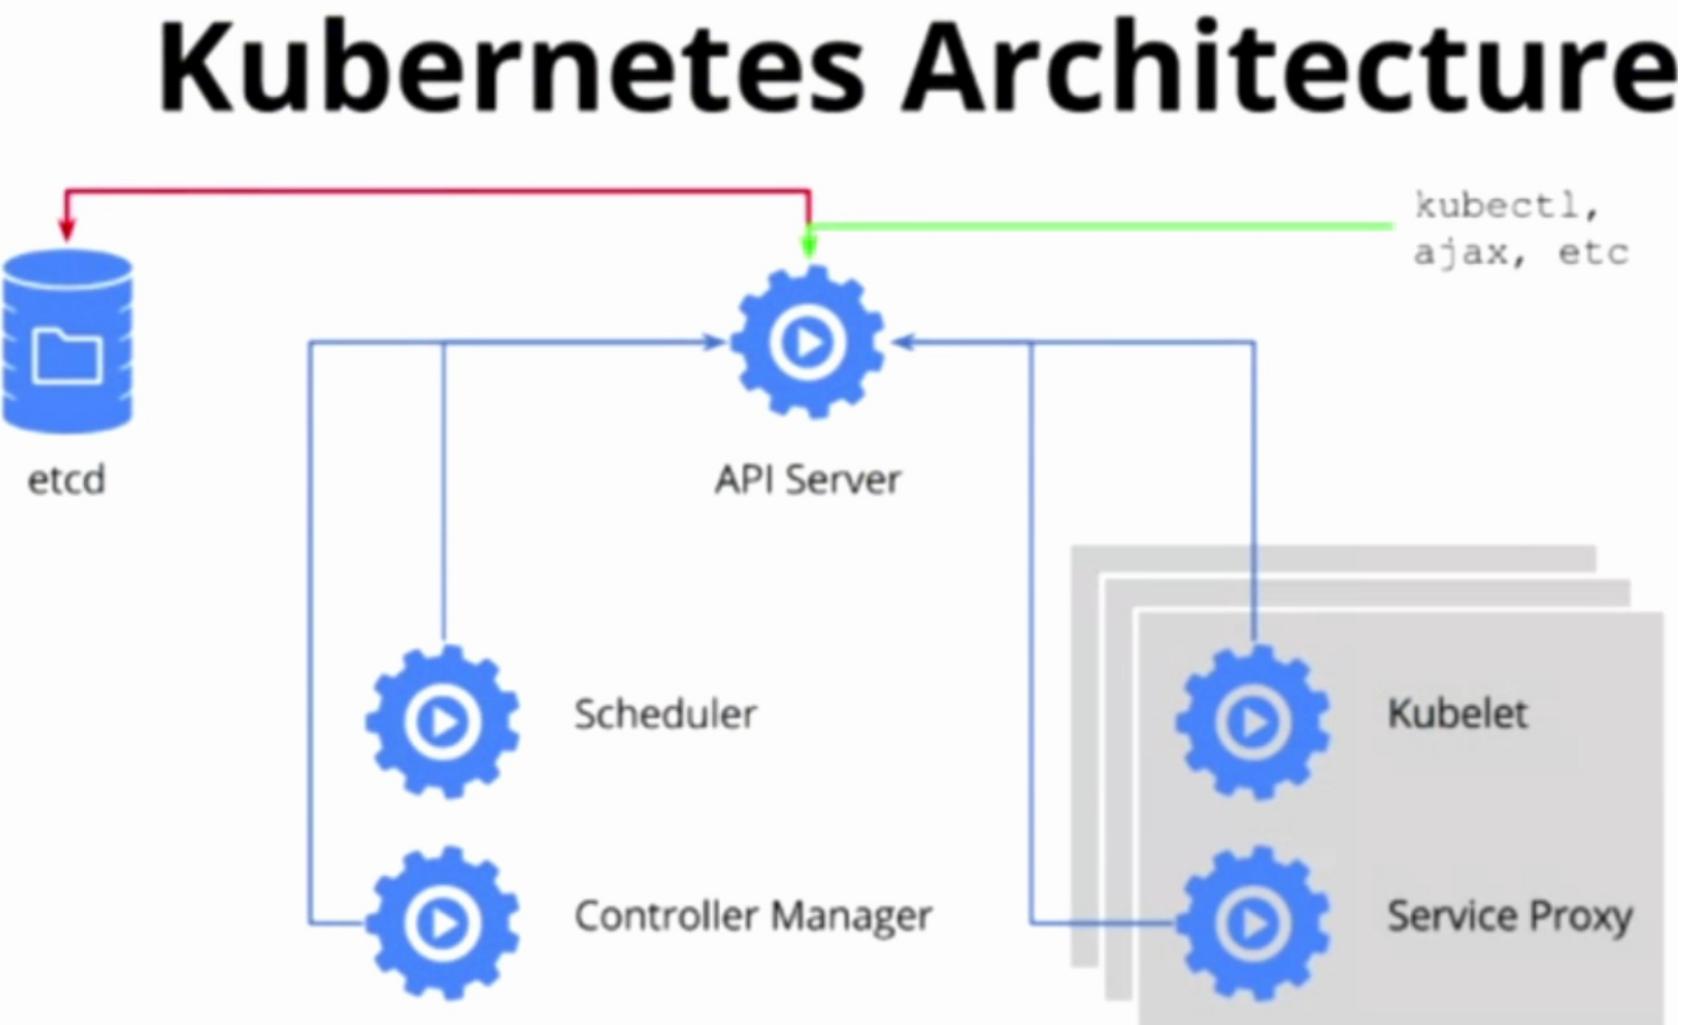
Integrated in:

- Stackanetes, Mirantis
- RedHat OpenShift PaaS
- Deis Paas
- EBay : Kubernetes + OVS
- CoreOS: Tectonic (commercial Kubernetes offering)



Architecture

Kubernetes





Concepts

- Cluster

Kubernetes^{Node}

- Pod
- Replication controller
- Service
- Label



Using Kubernetes

Step1

Kubernetes^{xxxx}

Step2



Getting started

An excellent place to start is with the following tutorials

Kubernetes
minikube?
...
...

Industry Players

Players

Choices made by Industry Players

- RedHat: Completely redesigned their OpenShift PaaS to use Docker Containers and Kubernetes, and created Project Atomic
- CoreOS: CoreOS, created the company 6 months after Docker was announced with a goal of providing **GIFFE**
- Google: Kubernetes used for GCP
- MicroSoft: Committed to port Docker to Windows (Windows Server 2016, Azure)
- VMWare

So isn't it time we told you what to
choose?

So isn't it time we told you what to
choose?

... well we'll provide some
guidelines at least ...

What's common

They are converging on many points

They are tending to add 'declarative specification' capabilities.

It's no longer feasible for an operator to decide on which node to deploy especially when complex constraints exist

- making use of specialized hardware, e.g. SSD best for some operations
- adapting to hardware failures

An operator specifies the "desired state" and the orchestrator does the rest.

What's different

Rancher lightweight

@hguemar @mjbright @mariolet

Comparison

Feature	Swarm	Kubernetes	Mesos
Declarative		Yes	

Rancher lightweight

@hguemar @mjbright @mariolet

Hands on ...

@hguemar @mjbright @mariolet

Hands-on

Come along

**This afternoon's tutorial session led by Mario:
Tuesday, October 4 - 15:30 - 16:20**

5 Containers for 5 Languages: Patterns for Software Development Using Containers - Mario Loriedo, Red Hat

**Tomorrow's lab session led by Haikel:
Wednesday, October 5 - 11:00 - 12:50**

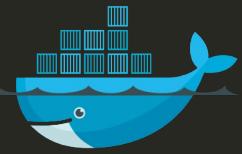
Container Orchestration Lab: Swarm, Mesos, Kubernetes - Haïkel Guémar, Fedora Project

Lab setup instructions [here](#)

- Docker Swarm
- Kubernetes
- Apache Mesos

@hguemar @mjbright @mariolet

Docker Swarm Demo



Questions?

Thank you

@hguemar @mjbright @mariolet

Resources

@hguemar @mjbright @mariolet

Resources Books

Publisher	Title	Author
O'Reilly	Docker Cookbook	Sébastien Goasguen
O'Reilly	Docker Up & Running	Karl Matthias, Sean P. Kane
O'Reilly	Using Docker [Early Access]	Adrian Mouat
O'Reilly	Kubernetes Up & Running	Kelsey Hightower
Manning	[MEAP] CoreOS in Action	Matt Bailey
Manning	[MEAP] Kubernetes in Action	Marko Lukša

@hguemar @mjbright @mariolet

Resources Articles/Organisms

Cloud Native Computing Foundation - Kubernetes,
Prometheus <https://cncf.io/>

"Kubernetes the Hard Way, Kelsey Hightower" -
<https://github.com/kelseyhightower/kubernetes-the-hard-way> *"Kubernetes User Guide, Walkthrough"* -
<http://kubernetes.io/docs/user-guide/walkthrough/>

Resources Videos

- June 2016 - Container Orchestration Wars, Karl Isenberg, Mesosphere
- Mar 2016 - Container Orchestration with Kubernetes, Docker Swarm & Mesos-Marathon - Adrian Mouat, Container Solutions
- Jan 2016 - Docker, Kubernetes, and Mesos: Compared.,,Adrian Otto, Southern California Linux Expo

Repos

@hguemar @mjbright @mariolet



Documentation

Kubernetes

- **Getting started guides**
 - [Creating a Kubernetes Cluster](#)
 - port Kubernetes to a new environment
 - in [Getting Started from Scratch](#)
- **User documentation**
 - to run programs on an existing Kubernetes cluster
 - [Kubernetes User Guide: Managing Applications](#)
 - the [Kubectl Command Line Interface](#) is a detailed reference on the `kubectl` CLI
 - [User FAQ](#)



Documentation - 2

Kubernetes

- **Cluster administrator documentation**
 - for people who want to create a Kubernetes cluster and administer it
 - in the [Kubernetes Cluster Admin Guide](#)
- **Developer and API documentation**
 - to write programs using the Kubernetes API, write plugins or extensions, or modify core code
 - [Kubernetes Developer Guide](#)
 - [notes on the API](#)
 - [API object documentation](#), a detailed description of all fields found in the core API objects
- **Walkthroughs and examples**
 - hands-on introduction and example config files
 - in the [user guide](#)
 - in the [docs/examples directory](#)
- **Contributions from the Kubernetes community**
 - in the [docs/contrib directory](#)



Documentation 3

Kubernetes

- **Design documentation and design proposals**
 - to understand the design of Kubernetes, and feature proposals
 - [Kubernetes Design Overview](#) and the [docs/design directory](#)
 - [docs/proposals directory](#)
- **Wiki/FAQ**
 - the [wiki](#)
 - [troubleshooting guide](#)

Community, discussion, contribution, and support

Consider joining the [Cloud Native Computing Foundation](#). For details about who's involved and how Kubernetes plays a role, read [their announcement](#).